

Promptly Data Pipeline Report

1. Overview

The **Promptly Data Pipeline** is a structured workflow designed to process **user queries and documents (PDF, TXT files)** for an **AI-driven document-based Q&A system**. The system uses **Retrieval-Augmented Generation (RAG)** to fetch relevant answers based on user queries.

This pipeline is built using **Apache Airflow** for orchestration, **Supabase** for database management, **Google Cloud Storage (GCS)** for document storage, and **Data Version Control (DVC)** for dataset tracking. The pipeline follows **MLOps best practices** to ensure **data consistency, versioning, and automation**.

This report details all pipeline components, architecture, and justifications for any sections that are not applicable.

2. Key Components in the Data Pipeline

The data pipeline follows a structured approach, covering **data acquisition, preprocessing, versioning, orchestration, tracking, and logging**.

2.1 Data Acquisition

Data Sources:

- **User Queries:** Retrieved from the `conversations` table in **Supabase**.
- **Documents:** Uploaded by users and stored in **Google Cloud Storage (GCS)**.

Implementation:

- Queries are fetched using the `supabase_utils.py` script.
- Documents are retrieved from the `promptly-chat` bucket in GCS.
- All dependencies and external configurations are managed in `requirements.txt`.

Reproducibility:

- Data retrieval steps are automated through Airflow DAGs.
- Dependencies and configurations are versioned using **DVC** to ensure replicability.

2.2 Data Preprocessing

Query Processing:

- **Schema validation** using `validate_schema.py` in both `rag` and `user queries pipeline`..
- **Text cleaning** using `data_utils.py` (stopword removal, special character removal, lowercasing).
- **Tokenization and Lemmatization** using `NLTK` and `WordNetLemmatizer`.

Document Processing:

- **Personally Identifiable Information (PII) Detection and Redaction** using `Presidio-based Named Entity Recognition (NER)` in `check_pii_data.py`.
- **Text Chunking** using `MarkdownSplitter` in `rag_utils.py`.

Output:

- Preprocessed **user queries and documents** are stored in **Supabase and GCS**.
-

2.3 Test Modules

Unit Testing Frameworks:

- `pytest`
- `unittest`

Implemented Test Cases:

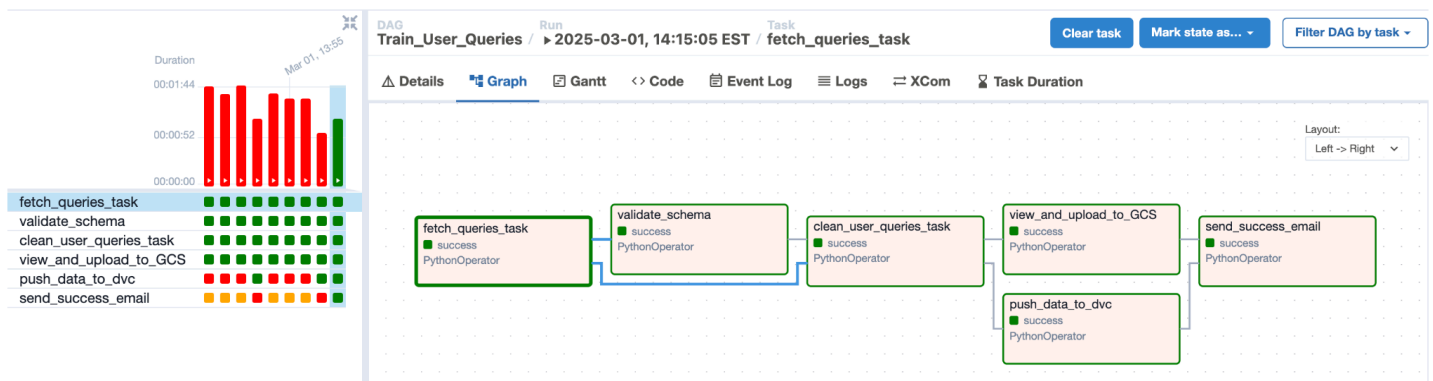
1. **PII Detection and Redaction** (`test_data_pii_redact_test.py`)
 2. **RAG Pipeline** (`test_rag_pipeline.py`)
 3. **User Query Pipeline** (`test_user_queries.py`)
-

2.4 Pipeline Orchestration (Airflow DAGs)

User Queries DAG: Train_User_Queries

Workflow:

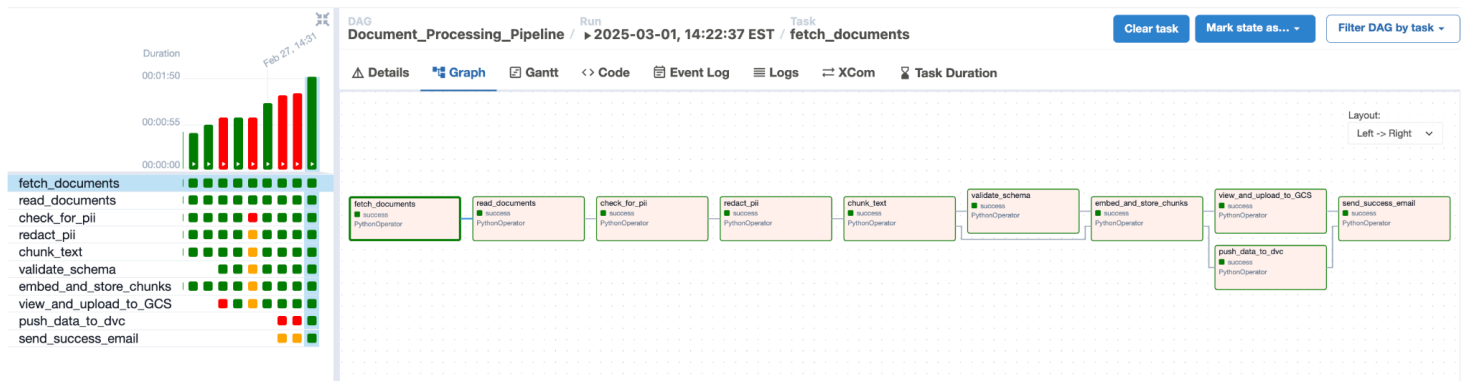
1. `fetch_queries_task`: Retrieves queries from Supabase.
2. `validate_schema`: Ensures schema consistency.
3. `clean_user_queries_task`: Cleans and preprocesses text.
4. `view_and_upload_to_GCS`: Saves processed queries to **GCS**.
5. `push_data_to_dvc`: Enables versioning via **DVC**.
6. `send_success_email`: Sends completion notifications.



Document Processing DAG: Document_Processing_Pipeline

Workflow:

1. `fetch_documents`: Retrieves uploaded PDFs/TXT files.
2. `read_documents`: Extracts text using `pymupdf4llm`.
3. `check_for_pii`: Identifies PII using **Presidio**.
4. `redact_pii`: Redacts sensitive data.
5. `chunk_text`: Splits documents into structured chunks.
6. `validate_schema`: Ensures correct formatting.
7. `embed_and_store_chunks`: Generates embeddings using **Nomic** and stores them in **Supabase**.
8. `view_and_upload_to_GCS`: Stores chunked documents in **GCS**.
9. `push_data_to_dvc`: Enables versioning via **DVC**.
10. `send_success_email`: Sends completion notifications.



2.5 Data Versioning with DVC

Implementation:

- **User Queries:** Versioned using **DVC** (`preprocessed_user_queries.csv.dvc`).
- **Processed Documents:** Versioned using **DVC** (`preprocessed_docs_chunks.csv`).
- **GCS Remote Storage:** Configured in `.dvc/config`.

DVC Workflow:

```
dvc init
```

```
dvc remote add gcs_remote gs://promptly-chat
```

```
dvc push
```

Benefits of Using DVC:

- Ensures **reproducibility** for dataset modifications.
- Allows **rollback** to previous versions.
- Provides **data lineage tracking** for query processing and document embedding versions.

2.6 Tracking and Logging

Logging Frameworks Used:

- **Airflow DAG Logs**
- **Python Logging Module**
- **Supabase Database Logging**

Implementation Example:

```
import logging

logging.basicConfig(level=logging.INFO)
```

Error Handling Strategies:

- **Retries in Airflow DAGs**
 - **Failure Notifications via Email**
 - **Exception Handling in Scripts**
-

2.7 Data Schema and Statistics Generation

Implementation:

- **Schema Validation:** Implemented in `validate_schema.py` before storing data in Supabase.
- **Statistics Generation:** Supabase logs track query/document statistics.

Justification for Not Using TFDV or MLMD:

- The system is **not an ML training pipeline**; it is focused on **retrieval**.
 - Instead of **TensorFlow Data Validation (TFDV)**, schema checks are **embedded in preprocessing scripts**.
-

2.8 Anomaly Detection and Alerts

Implemented Anomaly Detection:

- **Missing Data Checks:** Handled in `validate_schema.py`.
- **Unexpected Formats Detection:** Managed in schema as well as in `data_utils.py`.

Anomaly Alerts / Notifications:

- The system operates **on batch processing**, not **real-time streaming data**.
 - Pipeline failures are monitored through **Airflow Logs** and **email notifications**.
-

3. Data Bias Detection Using Data Slicing

Justification for Not Implementing Bias Detection:

- Our chatbot processes IT-related documents stored in a Retrieval-Augmented Generation (RAG) system, which contains purely technical content free from demographic or subjective influences.
 - No Demographic Features: Our data lacks attributes like age, gender, or location, making subgroup analysis irrelevant.
 - No Subjective Bias: The dataset consists of factual IT documentation, ensuring objectivity in responses.
 - Unnecessary Overhead: Bias detection tools like TFMA or Fairlearn are designed for datasets with social implications. Applying them here would add complexity without meaningful insights.
 - Since bias is not an inherent risk in our dataset, we will be focusing on enhancing precision and relevance in chatbot responses.
-

4. Folder Structure and Code Organization

The project follows **MLOps best practices** for modularity and scalability.

```
├─ assets/
|   └─ process_user_queries_dag.png  # User Query Pipeline DAG
|   └─ rag_data_pipeline_dag.png    # Data Pipeline Workflow Diagram
|
├─ data_pipeline/
|   └─ dags/
|       └─ dataPipeline.py  # User Queries DAG
```

```
| | └─ rag_data_pipeline.py # Document Processing DAG
| | └─ scripts/
| | | └─ email_utils.py # Email notifications
| | | └─ upload_data_GCS.py # GCS Uploading
| | | └─ data_preprocessing/
| | | | └─ check_pii_data.py # PII Detection
| | | | └─ validate_schema.py # Schema Validation
| | | | └─ data_utils.py # Query Cleaning Functions
| | | └─ supadb/
| | | | └─ supabase_utils.py # Supabase Integration
| | | └─ rag/
| | | | └─ rag_utils.py # Chunking & Embeddings
| | └─ tests/
| | | └─ test_data_pii_redact.py # Tests for PII detection
| | | └─ test_rag_pipeline.py # Tests for RAG doc pipeline
| | | └─ test_user_queries.py # Tests for User Queries pipeline
| └─ config.py # API Keys & Configurations
| └─ README.md # Project Documentation
|
└─ data/
    └─ rag_documents/ # Original PDFs & Text Files
    └─ preprocessed_docs_chunks.csv # Cleaned & Chunked Data
    └─ preprocessed_user_data.csv # Processed Queries
|
```

```
|— .dvc/ # DVC Configuration
|— .gitignore
|— requirements.txt # Dependencies
```

5. Replicability and Reproducibility

Steps to Set Up:

1. Clone the Repository:

```
git clone https://github.com/RajivShah1798/promptly.git
cd promptly-data-pipeline
```

2. Install Dependencies:

```
pip install -r requirements.txt

Setup env file and add SUPABASE_KEY and URL
```

3. Initialize DVC:

```
dvc init

dvc remote add gcs_remote gs://promptly-chat

dvc pull
```

4. Run Airflow DAGs:

```
airflow db init

airflow scheduler & airflow webserver

airflow dags trigger Train_User_Queries
```

6. Evaluation Criteria

The pipeline meets all **required evaluation criteria**:

- **Proper Documentation**
- **Modular Code & Error Handling**
- **DVC Versioning**
- **Pipeline Optimization**
- **Logging and Alerts**